

Vishal Talwar

Data Science And Business Analytics Intern @Spark Foundation(TSF)

Simple Linear Regression

```
In [116]: import pandas as pd      #For Data Manipulation
import numpy as np      #For Working with n-dimensional arrays

import matplotlib.pyplot as plt      #For visualization of data
import seaborn as sns      #For visulatization purpose

from sklearn.model_selection import train_test_split as ts #For splitting our data
from sklearn.linear_model import LinearRegression #algorithm for making model
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
```

```
In [3]: df =pd.read_csv("https://raw.githubusercontent.com/AdiPersonalWorks/Random/master/
```

```
In [5]: #For inspecting first five rows of data
df.head()
```

Out[5]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [6]: #For Inspecting rows selected at random  
df.sample(10)
```

Out[6]:

	Hours	Scores
21	4.8	54
2	3.2	27
22	3.8	35
13	3.3	42
20	2.7	30
1	5.1	47
18	6.1	67
12	4.5	41
3	8.5	75
24	7.8	86

```
In [9]: #Checking Shape of our data  
print("No. of rows are %i"%df.shape[0])  
print("No of columns are %i "%df.shape[1])
```

No. of rows are 25
No of columns are 2

```
In [10]: #Checking data types of columns  
df.dtypes
```

Out[10]: Hours float64
Scores int64
dtype: object

```
In [11]: #Checking Null values in our dataset  
df.isnull().sum()
```

Out[11]: Hours 0
Scores 0
dtype: int64

```
In [17]: #Checking for duplicate rows in our data  
df.duplicated().sum()
```

Out[17]: 0

```
In [19]: #Looking at the statistical property of columns  
df.describe()
```

Out[19]:


	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
In [21]: #Checking for correlation in our data to check linearity assumption in Linear Reg  
df.corr()
```


Out[21]:

	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

```
In [22]: #Univariate Analysis  
hour=sns.distplot(df['Hours'])
```




```
In [24]: Score=sns.distplot(df["Scores"])
```




```
In [39]: Hours=sns.boxplot(df['Hours'])
```

#No outlier




```
In [34]: Score=sns.boxplot(df["Scores"])
```

#No outliers



```
In [33]: #Bivariate Analysis  
sns.scatterplot(df["Hours"],df["Scores"]);  
plt.title("Hours studied vs Marks Scored");
```

#Data is positive correlated



```
In [68]: #Assigning column to dependent and independent variable
X=df['Hours'].values.reshape(-1,1)
y=df['Scores'].values
```

```
In [69]: #Splitting our data for training and predicting purpose
X_train,X_test,y_train,y_test=ts(X,y,test_size=0.2)
```

```
In [70]: print(X_train.shape)
print(X_test.shape)
```

```
(20, 1)
(5, 1)
```

```
In [71]: print(y_train.shape)
print(y_test.shape)
```

```
(20,)
(5,)
```

```
In [77]: #Training our model
model=LinearRegression()
model.fit(X_train,y_train)
```

```
Out[77]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [88]: print(model.coef_)
print(model.intercept_)
```

```
#Interpreting Model Coefficient
```

```
#A"Unit" increase in number of hours studied is associated with 9.955"units" in r
```

```
[9.95084503]
1.4847149599514253
```

```
In [89]: #Making prediction using test data
y_pred=model.predict(X_test)
```

```
In [90]: #Comparing actual vs predicted scores
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
df
```

```
Out[90]:
```

	Actual	Predicted
0	30	26.361828
1	67	62.184870
2	81	84.076729
3	69	75.120968
4	60	56.214363

```
In [118]: #Evaluating our model
print("Mean absolute error is ",mean_absolute_error(y_pred,y_test))
print("r square value is",r2_score(y_pred,y_test))
```

Mean absolute error is 4.2873274190705635
r square value is 0.9500542668240033

```
In [119]: #Conclusion
p=model.predict([[9.25]])
print("Number of hours studied are %f and predicted score is %f"%(9.25,p))
```

Number of hours studied are 9.250000 and predicted score is 93.530032

In []: